



# COOPERATIVE GROUPS

## FLEXIBLE GROUPS OF THREADS

21 June 2023 | Andreas Herten | Forschungszentrum Jülich *Handout Version*

# Overview, Outline

## At a Glance

- Cooperative Groups: New model to work with thread groups
- Thread groups are entities, intrinsic function as member functions

## Contents

Motivation

Basis

Cooperative Groups

Introduction

Thread Groups Overview

Thread Blocks

Task 1

Tiling Groups

Dynamic Size

Static Size

Coalesced Groups

Binary Partition

Labeled Partition

Larger Groups

Task 2

Warp-Synchronous Programming

Overview

Task 3

Collective Operations

Block Clusters

Conclusions

# Gather Last-Minute Material

Now run

```
jsc-material-reset-03
```


Place cursor in box when done:

I'm done!

# Standard CUDA Threading Model

## Before CUDA 9

- Many threads, combined into blocks, on a grid; in 3D
- Operation: Single Instruction, Multiple Threads (SIMT)
- Thread waiting for result of instruction? Use computational resource with other threads in meantime!
- Group of threads execute in lockstep: **Warp** (currently 32 threads)
  - Same instructions
  - Branching possible
  - Predicates (and masks)
- Shared memory: Fast, shared between threads of block
- Synchronization between threads of blocks:  
`__syncthreads()` – barrier for all threads of block



# Cooperative Groups

## Introduction

# New Model: Cooperative Groups

- Motivation to extend classical model

**Algorithmic** Not all algorithms map easily to available synchronization methods;  
**synchronization** should be more flexible; easier to utilize low-level concepts

**Design** Make groups of threads explicit **entities**


**Hardware** Access new **hardware features** (*Independent Thread Scheduling, Thread Block Clusters*)

## → Cooperative Groups (CG)

*A flexible model for synchronization and communication within groups of threads.*

- All in **namespace cooperative\_groups** (cooperative\_groups.h header)
- Following in text: `cooperative_groups::func()` → `cg::func()`  
**namespace cg = cooperative\_groups;**

# Division of Thread Blocks



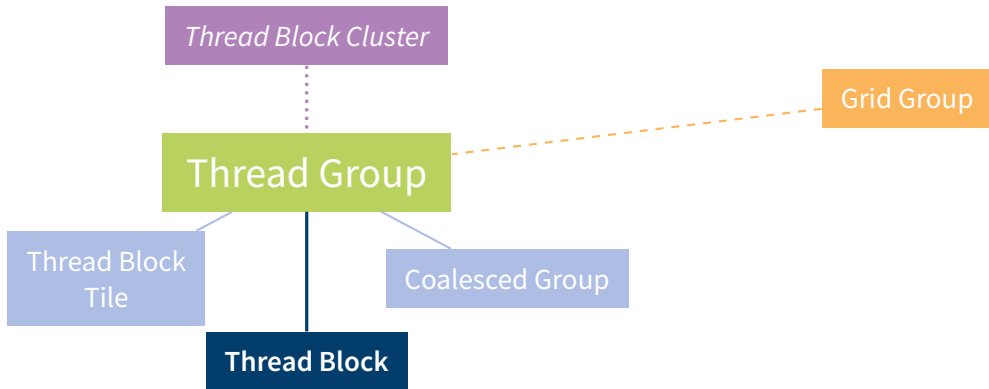
- Start with block of certain size
  - Divide into smaller sub-groups
  - Continue diving, if algorithm makes it necessary
  - Methods for dynamic or static divisions (*tiles*)
  - In each level: thread of group has unique ID (local index instead of global index)
- Use functions and collectives on sub-set of all threads

# Cooperative Groups

## Thread Groups Overview



# Thread Group Landscape




# Common Methods of Cooperative Groups

- Fundamental type: `thread_group`
- Every CG has following member functions
  - `sync()` Synchronize the threads of this group (alternative `cg::sync(g)`)  
*Before: `__syncthreads()` for whole block*
  - `thread_rank()` Get unique ID of current thread in this group (*local index*)  
*Before: `threadIdx.x` for index in block*
  - `size()` Number of threads in this group  
*Before: `blockDim.x` for number of threads in block*
  - `is_valid()` *Group is technically ok*

# Cooperative Groups

## Thread Blocks

# Cooperative Thread Blocks



- Easiest entry point to thread groups: `cg::this_thread_block()`

- Additional member functions

`thread_index()` Thread index within block (3D)

`group_index()` Block index within grid (3D)

- Blocks (and groups) are now concrete entities

→ Design functions to represent this!

# Example: Print Rank Function

```
__device__ void printRank(cg::thread_group g) {  
    printf("Rank %d\n", g.thread_rank());  
}  
  
__global__ void allPrint() {  
    cg::thread_block b = cg::this_thread_block();  
  
    printRank(b);  
}  
  
int main() {  
    allPrint<<<1, 23>>>();  
}
```

# Task Base Code: Shared Memory Reduction

Inner logic: Function

```
int * array;  
cudaMallocManaged(&array, sizeof(int) * N);
```

```
for (int i = 0; i < N; i++)  
    array[i] = rand() % 1024;
```

```
int blocks = 1;  
int threads = N;  
maxKernel<<<blocks, threads, threads * sizeof(int)>>>(array);
```

*Allocate this much shared memory per block*

```
__global__ void maxKernel(int * array) {  
    extern __shared__ int shmem_temp[]; // threads * sizeof(int)
```

```
int threadIndex = threadIdx.x;  
int myValue = array[threadIndex];
```

One value for each thread

```
int maxValue = maxFunction(shmem_temp, myValue);
```

# Task Base Code: Shared Memory Reduction

Inner logic: Function

```
int * array;
cudaMallocManaged(&array, sizeof(int) * N);

for (int i = 0; i < N; i++)
    array[i] = rand() % 1024;

int blocks = 1;
int threads = N;
maxKernel<<<blocks, threads, threads * sizeof(int)>>>(array);

Allocate this much shared memory per block
__global__ void maxKernel(int * array) {
    extern __shared__ int shmem_temp[]; // threads * sizeof(int)

    int threadIndex = threadIdx.x;
    int myValue = array[threadIndex];

    int maxValue = maxFunction(shmem_temp, myValue);
```

One value for each thread



**JÜLICH**  
Forschungszentrum

JÜLICH  
SUPERCOMPUTING  
CENTRE

Call function with

# Task Base Code: Shared Memory Reduction

Inner logic: Function

```
int * array;
cudaMallocManaged(&array, sizeof(int) * N);

for (int i = 0; i < N; i++)
    array[i] = rand() % 1024;

int blocks = 1;
int threads = N;
maxKernel<<<blocks, threads, threads * sizeof(int)>>>(array);

Allocate this much shared memory per block
__global__ void maxKernel(int * array) {
    extern __shared__ int shmem_temp[]; // threads * sizeof(int)

    int threadIndex = threadIdx.x;
    int myValue = array[threadIndex];

    int maxValue = maxFunction(shmem_temp, myValue);
```

One value for each thread



**JÜLICH**  
Forschungszentrum

JÜLICH  
SUPERCOMPUTING  
CENTRE

Call function with



# Task Base Code: Shared Memory Reduction

Inner logic: Function

```
int * array;
cudaMallocManaged(&array, sizeof(int) * N);

for (int i = 0; i < N; i++)
    array[i] = rand() % 1024;

int blocks = 1;
int threads = N;
maxKernel<<<blocks, threads, threads * sizeof(int)>>>(array);

Allocate this much shared memory per block
__global__ void maxKernel(int * array) {
    extern __shared__ int shmem_temp[]; // threads * sizeof(int)

    int threadIndex = threadIdx.x;
    int myValue = array[threadIndex];

    int maxValue = maxFunction(shmem_temp, myValue);
```

One value for each thread



**JÜLICH**  
Forschungszentrum

JÜLICH  
SUPERCOMPUTING  
CENTRE

Call function with

# Task Base Code: Shared Memory Reduction

Inner logic: Function

```
int * array;
cudaMallocManaged(&array, sizeof(int) * N);

for (int i = 0; i < N; i++)
    array[i] = rand() % 1024;

int blocks = 1;
int threads = N;
maxKernel<<<blocks, threads, threads * sizeof(int)>>>(array);

Allocate this much shared memory per block
__global__ void maxKernel(int * array) {
    extern __shared__ int shmem_temp[]; // threads * sizeof(int)

    int threadIndex = threadIdx.x;
    int myValue = array[threadIndex];

    int maxValue = maxFunction(shmem_temp, myValue);
```

One value for each thread



**JÜLICH**  
Forschungszentrum

JÜLICH  
SUPERCOMPUTING  
CENTRE

Call function with

# Task Base Code: Shared Memory Reduction

Inner logic: Function

```
int * array;
cudaMallocManaged(&array, sizeof(int) * N);

for (int i = 0; i < N; i++)
    array[i] = rand() % 1024;

int blocks = 1;
int threads = N;
maxKernel<<<blocks, threads, threads * sizeof(int)>>>(array);

Allocate this much shared memory per block
__global__ void maxKernel(int * array) {
    extern __shared__ int shmem_temp[]; // threads * sizeof(int)

    int threadIndex = threadIdx.x;
    int myValue = array[threadIndex];

    int maxValue = maxFunction(shmem_temp, myValue);
```

One value for each thread



**JÜLICH**  
Forschungszentrum

JÜLICH  
SUPERCOMPUTING  
CENTRE

Call function with

# Task Base Code: Shared Memory Reduction

Inner logic: Function

```
int * array;
cudaMallocManaged(&array, sizeof(int) * N);

for (int i = 0; i < N; i++)
    array[i] = rand() % 1024;

int blocks = 1;
int threads = N;
maxKernel<<<blocks, threads, threads * sizeof(int)>>>(array);

Allocate this much shared memory per block
__global__ void maxKernel(int * array) {
    extern __shared__ int shmem_temp[]; // threads * sizeof(int)

    int threadIndex = threadIdx.x;
    int myValue = array[threadIndex];

    int maxValue = maxFunction(shmem_temp, myValue);
```

One value for each thread



**JÜLICH**  
Forschungszentrum

JÜLICH  
SUPERCOMPUTING  
CENTRE

Call function with

# Task Base Code: Shared Memory Reduction

Inner logic: Function

```
int * array;  
cudaMallocManaged(&array, sizeof(int) * N);
```

```
for (int i = 0; i < N; i++)  
    array[i] = rand() % 1024;
```

```
int blocks = 1;  
int threads = N;  
maxKernel<<<blocks, threads, threads * sizeof(int)>>>(array);
```

*Allocate this much shared memory per block*

```
__global__ void maxKernel(int * array) {  
    extern __shared__ int shmem_temp[]; // threads * sizeof(int)
```

```
int threadIndex = threadIdx.x;  
int myValue = array[threadIndex];
```

One value for each thread

```
int maxValue = maxFunction(shmem_temp, myValue);
```

# Task Base Code: Shared Memory Reduction

Inner logic: Function

```
int * array;  
cudaMallocManaged(&array, sizeof(int) * N);
```

```
for (int i = 0; i < N; i++)  
    array[i] = rand() % 1024;
```

```
int blocks = 1;  
int threads = N;  
maxKernel<<<blocks, threads, threads * sizeof(int)>>>(array);
```

*Allocate this much shared memory per block*

```
__global__ void maxKernel(int * array) {  
    extern __shared__ int shmem_temp[]; // threads * sizeof(int)
```

```
int threadIndex = threadIdx.x;  
int myValue = array[threadIndex];
```

One value for each thread

```
int maxValue = maxFunction(shmem_temp, myValue);
```

# Task Base Code: Shared Memory Reduction

Inner logic: Find the max

```
int * array;  
cudaMallocManaged(&array, sizeof(int) * N);
```

```
for (int i = 0; i < N; i++)  
    array[i] = rand() % 1024;
```

```
int blocks = 1;  
int threads = N;  
maxKernel<<<blocks, threads, threads * sizeof(int)>>>(array);
```

```
Allocate this much shared memory per block  
__global__ void maxKernel(int * array) {  
    extern __shared__ int shmem_temp[]; // threads * sizeof(int)
```

```
    int threadIndex = threadIdx.x;  
    int myValue = array[threadIndex];
```

```
    int maxValue = maxFunction(shmem_temp, myValue);
```

One value for each thread



**JÜLICH**  
Forschungszentrum

JÜLICH  
SUPERCOMPUTING  
CENTRE

Call function with

# Task Base Code: Shared Memory Reduction

Inner logic:

```
int * array;
```

```
cudaMallocManaged(&array, sizeof(int) * N);
```

```
for (int i = 0; i < N; i++)  
    array[i] = rand() % 1024;
```

```
int blocks = 1;  
int threads = N;  
maxKernel<<<blocks, threads, threads * sizeof(int)>>>(array);
```

```
__global__ void maxKernel(int * array) {  
    extern __shared__ int shmem_temp[]; // threads * sizeof(int)
```

```
    int threadIndex = threadIdx.x;  
    int myValue = array[threadIndex];
```





# Task Base Code: Shared Memory Reduction

Inner logic: Find the max

```
int * array;
```

```
cudaMallocManaged(&array, sizeof(int) * N);
```

```
for (int i = 0; i < N; i++)
```

```
array[i] = rand() % 1024;
```

```
int blocks = 1;
```

```
int threads = N;
```

```
maxKernel<<<blocks, threads, threads * sizeof(int)>>>(array);
```


```
__global__ void maxKernel(int * array) {
```

```
    extern __shared__ int shmem_temp[]; // threads * sizeof(int)
```

```
    int threadIndex = threadIdx.x;
```

```
    int myValue = array[threadIndex];
```

```
    int maxValue = maxFunction(shmem_temp, myValue);
```



# Task Base Code: Shared Memory Reduction

Inner logic: Find the max

```
int * array;  
cudaMallocManaged(&array, sizeof(int) * N);
```

```
for (int i = 0; i < N; i++)  
    array[i] = rand() % 1024;
```

```
int block = 1024;  
int threads = 1024;  
maxKernel<<blocks, threads, threads * sizeof(int)>>>(array);
```

```
Allocate this much shared memory per block  
__global__ void maxKernel(int * array) {  
    extern __shared__ int shmem_temp[]; // threads * sizeof(int)
```

```
    int threadIndex = threadIdx.x;  
    int myValue = array[threadIndex];
```

```
    int maxValue = maxFunction(shmem_temp, myValue);
```



One value for each thread



**JÜLICH**  
Forschungszentrum

JÜLICH  
SUPERCOMPUTING  
CENTRE

# Task Base Code: Shared Memory Reduction

Inner logic: For each thread

```
int * array;  
cudaMallocManaged(&array, sizeof(int) * N);
```

```
for (int i = 0; i < N; i++)  
    array[i] = rand() % 1024;
```

```
int blocks = 1;  
int threads = N;
```

```
maxKernel(<blocks, threads, threads * sizeof(int)>>>(array);
```

```
__global__ void maxKernel(int * array;
```

```
extern __shared__ int shmem_temp[]; // threads * sizeof(int)
```

```
int threadIndex = threadIdx.x;
```

```
int myValue = array[threadIndex];
```

```
int maxValue = maxFunction(shmem_temp, myValue);
```



# Task Base Code: Shared Memory Reduction

Inner logic: Find the max

```
int * array;  
cudaMallocManaged(&array, sizeof(int) * N);
```

```
for (int i = 0; i < N; i++)  
    array[i] = rand() % 1024;
```

```
int blocks = 1;  
int threads = N;  
maxKernel<<blocks, threads, threads * sizeof(int)>>>(array);
```

*Allocate this much shared memory per block*

```
__global__ void maxKernel(int * array) {
```

```
    extern __shared__ int shmem_temp[]; // threads * sizeof(int)
```

```
    int threadIndex = threadIdx.x;  
    int myValue = array[threadIndex];
```

```
    int maxValue = maxFunction(shmem_temp, myValue);
```

```
    int callFunctionWith
```

# Task Base Code: Shared Memory Reduction

Inner logic: Find the max

```
int * array;  
cudaMallocManaged(&array, sizeof(int) * N);
```

```
for (int i = 0; i < N; i++)  
    array[i] = rand() % 1024;
```

```
int blocks = 1;  
int threads = N;  
maxKernel<<<blocks, threads, threads * sizeof(int)>>>(array);
```

```
Allocate this much shared memory per block  
__global__ void maxKernel(int * array) {  
    extern __shared__ int shmem_temp[]; // threads * sizeof(int)
```

```
    int threadIndex = threadIdx.x;  
    int myValue = array[threadIndex];
```

```
    int maxValue = maxFunction(shmem_temp, myValue);
```

One value for each thread



**JÜLICH**  
Forschungszentrum

JÜLICH  
SUPERCOMPUTING  
CENTRE

Call function with

# Implementing a Cooperative Groups Kernel

TASK 1

From old to new


- Location of code: `03-Cooperative_Groups/exercises/tasks/task1`
- See `Instructions.md` for explanations
- Follow TODOs to port kernel/device function from traditional CUDA threading model to new CG model
- Compile with `make`, submit to batch system with `make run`
- See also [CUDA C programming guide](#) for details on Cooperative Groups

# Cooperative Groups

## Tiling Groups

# Tiles of Groups

## Dynamically-tiled



- Divide into smaller groups with `cg::tiled_partition()`
- Will automatically create smaller groups from parent group
- Examples
  - Create groups of size 32 of current block


```
cg::thread_group tile32 = cg::tiled_partition(cg::this_thread_block(), 32);
```
  - Create sub-groups of size 4

```
cg::thread_group tile4 = cg::tiled_partition(tile32, 4);
```
- **Note:** Currently, only supported partition sizes are 1, 2, 4, 8, 16, 32



# Tiles of Groups

## Statically-tiled: `thread_block_tile`




- Second version of function: `cg::tiled_partition<>()`
- Size of tile is template parameter
- Known at compile time! Optimizations possible!
- Partition size: 1, ..., 32, **64**, **128**, **256**, **512**! (*<A100: extra work needed*)
- Returns `thread_block_tile` object with additional member functions
  - `.shfl()`, `.shfl_down()`, `.shfl_up()`, `.shfl_xor()`
  - `.any()`, `.all()`, `.ballot()`; `.match_any()`, `.match_all()`
- Intrinsic functions to work with threads inside a warp (*more later*)
- Example

```
cg::thread_block_tile<32> tile32 = cg::tiled_partition<32>(cg::this_thread_block());
cg::thread_block_tile<4> tile4 = cg::tiled_partition<4>(tile32);
```

# Cooperative Groups

## Coalesced Groups

# Coalesced Group




- Get group of threads which is not diverged
- Threads have same state at point of API call
- `cg::coalesced_group active_threads = cg::coalesced_threads();`
- Example

```
cg::coalesced_group active_threads = cg::coalesced_threads();
if (i < 5) {
    cg::coalesced_group if_true_threads = cg::coalesced_threads();
    int rank = if_true_threads.thread_rank();
    cg::thread_group partition = cg::tiled_partition(if_true_threads, 2);
}
```

# Cooperative Groups

## Binary Partition

# Binary Partition




- Get group of coalesced threads for which a condition is either `true` or `false`
- Threads have same state at point of API call and belong to one of two *buckets*
- `cg::coalesced_group partitioned_threads = cg::binary_partition(group, condition);`
- *Beta* feature, details might change
- Example

```
cg::thread_block cta = cg::this_thread_block();
cg::thread_block_tile<32> tile32 = cg::tiled_partition<32>(cta);
auto subTile = cg::binary_partition(tile32, isEven(array[cta.thread_rank()])) ;
```

# Cooperative Groups

## Labeled Partition

# Labeled Partition



- Get group of coalesced threads for which a condition is equal
- Threads have same state at point of API call and belong to same *bucket*
- Extension of binary partition to general case
- `cg::coalesced_group partitioned_threads = cg::labeled_partition(group, condition);`
- *Beta* feature, details might change
- Example


```
cg::coalesced_group active = cg::coalesced_threads();  
auto labeledGroup = cg::labeled_partition(active, bucket);
```

# Cooperative Groups

## Larger Groups



# Grid Group



- Grid of blocks can also be entity now

- Synchronize across all blocks:

```
cg::grid_group grid = cg::this_grid();  
grid.sync();
```

- Condition

- 1 Blocks must be co-resident on device (Occupancy Calculator)
- 2 Kernel must be launched with Cooperative Launch API  
`cudaLaunchCooperativeKernel()` instead of `<<<, >>>` syntax

# Cooperative Groups with Tiled Partitions

## Sub-divisions

TASK 2

- Location of code: 03-Cooperative\_Groups/exercises/tasks/task2
- See Instructions.md for explanations
- Follow TODOs to tile a CG and use kernel from Task 1; **atomic operations** needed
- Compile with make, submit to batch system with make run
- See also [CUDA C programming guide](#) for details on Cooperative Groups

Aside!

# Aside: Atomic Operations


## Motivation

- Order execution of CUDA threads non-deterministic
- No problem, if each thread works on distinct data element
- What, if threads collaborate and share data? Read/Write to same element?

### → Atomic operations

- Safe way to read and write to memory position by different threads
- Data in global or shared memory
- Example: `atomicAdd(&array[i], myvalue)`
- See [CUDA Documentation](#)

`array[1] = array[1] + myvalue`



# Aside: Atomic Operations

## Examples

- First argument to function (always): address of a value to potentially change
- Old value of address usually returned
- `int atomicOp(int * removeVal, int myVal)`
- Examples

`atomicAdd(int* address, int val)` Add `val` to the value at address

`atomicExch(int* address, int val)` Store `val` at address location; return old value

`atomicMin(int* address, int val)` Store the minimum of `val` and the value at address at address location; return old value

`atomicCAS(int* address, int compare, int val)` The value at address is compared to `compare`. If true, `val` is stored at address; if false, the old value at address is stored. The old value at address is returned. Basic function: Compare And Swap

# Cooperative Groups with Tiled Partitions

## Sub-divisions

### TASK 2

- Location of code: `03-Cooperative_Groups/exercises/tasks/task2`
- See `Instructions.md` for explanations
- Follow TODOs to tile a CG and use kernel from Task 1; **atomic operations** needed
- Compile with `make`, submit to batch system with `make run`
- See also [CUDA C programming guide](#) for details on Cooperative Groups

# Warp-Synchronous Programming

# Warp-Level Intrinsic

- Smallest set of executed threads: Warp
- Warp: 32 threads executed in SIMT/SIMD fashion
- Exchange data between threads of warp
  - Global memory: Slow
  - Shared memory: Faster
  - Directly (registers): Even faster
- Safe method access without race conditions
  - Global/shared memory: Atomic operations
  - Registers: **Warp-aggregated Atomic operations**



# Warp Intrinsic Overview

- `shfl(int lane)` Copy data from a target warp lane; also: other flavors (next slide)
- `all(int pred)` If predicate (*comparison, relation*) evaluates to non-zero (*true*) for all threads, return non-zero (*true*)
- `any(int pred)` If predicate evaluates to non-zero for any thread, return non-zero
- `ballot(int pred)` Return a bit mask which has 1s set for all thread for which predicate evaluates to non-zero
- `match_any(T value)` Return a bit mask of threads which have same value of `value` as current thread; also: `match_all(T value)`
- Available as global device functions, with additional selection *mask* as first element (as `__shuf_l_sync()` etc.)
  - Available as **member functions** of a `cg::tiled_partition` group (as `g.shfl()` etc.)
  - Intrinsics automatically synchronize after operation – new since CUDA 9
  - Value can only be retrieved if targeted lane also invokes intrinsic
  - Per clock cycle: 32 shuffle instructions per SM → **very fast!**



# Warp Intrinsic Example

## Everyday I'm Shuffling

- `shfl()`: Copy data from target warp lane
- Different flavors
  - `shfl()` Copy data from warp lane with ID directly
  - `shfl_up()` Copy data from relative warp lane with lower ID (shuffle *upstream*)
  - `shfl_down()` Copy data from relative warp lane with higher ID (shuffle *downstream*)
  - `shfl_xor()` Copy data from relative warp lane with ID as calculated by a bitwise XOR
- **Example:** `shfl_down(value, N)` with  $N = 16, 8, \dots$

# Kernel → Warp-Level Reduction w/o Shared Memory

TASK 3

Expert level 11

- Location of code: `03-Cooperative_Groups/exercises/tasks/task3`
- See `Instructions.md` for explanations
- Follow TODOs to modify `maxKernel()` such that it uses warp-level atomic operations (and no shared memory)
- Compile with `make`, submit to batch system with `make run`
- See also [CUDA C programming guide](#) for details on warp-level functions

# Collective Operations

# Collective Operations

- In-group programming (ideally: warp-level programming) can get last bits of performance; but quite **advanced**
- Help: Collective operations on thread groups (**new** and slightly less **advanced**)
  - `cg::sync()` Synchronize threads in group
  - `cg::memcpy_async()` Copy from global to shared memory in group, non-blocking;  
also: `cg::wait`
  - `cg::reduce()` Reduction operation in group; hardware-accelerated operators:  
`plus()`, `less()`, `greater()`, `bit_and()`, `bit_xor()`, `bit_or()`
  - `cg::inclusive_scan()` Scan operation in group (also: `cg::exclusive_scan()`)

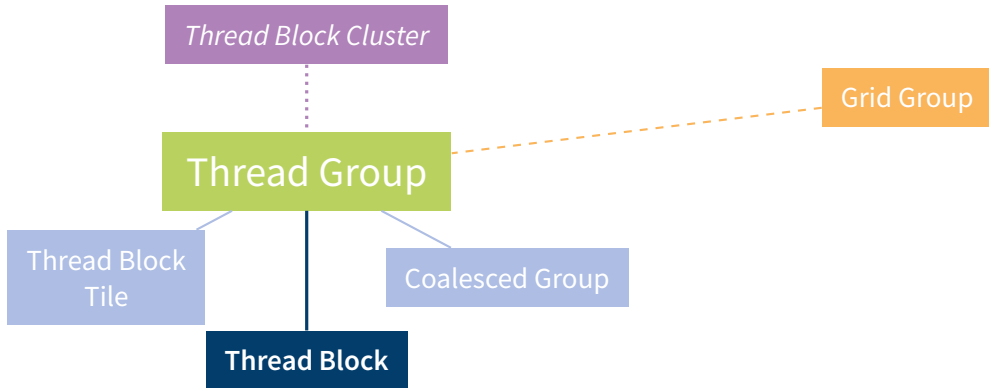
# Cooperative Reduce Collective Example

```
__shared__ int reduction_s[BLOCKSIZE];
cg::thread_block cta = cg::this_thread_block();
cg::thread_block_tile<32> tile = cg::tiled_partition<32>(cta);

const int tid = cta.thread_rank();
int value = A[tid];
reduction_s[tid] = cg::reduce(tile, value, cg::plus<int>());
// reduction_s contains tile-sum at all positions associated to tile
cg::sync(cta);
// Still to do: sum partial tile sums
```



# Block Clusters

# Thread Group Landscape



# Moare Hierarchy

- New feature available in next-gen H100 GPU in Compute Capability 9.0
- Extend hierarchy:  
Threads → Thread Blocks → Grids Threads → Thread Blocks → Thread Block Clusters → Grids
- Exposes the GPC (*GPU Processing Cluster*) hardware to software – **only through CG**
- Enables collaboration of some SMs of GPC; access shared memory (incl. atomics, like `sync( )`); max. 16 blocks per cluster





# Using Block Clusters

- Two possibilities for usage
  - 1 Through annotating intrinsic at kernel definition `__cluster_dims__(X,Y,Z)` (compile-time only)
  - 2 Through special kernel launch call `cudaLaunchKernelEx()` (also run-time)
- Guaranteed to be co-scheduled (running at same time)
- Use `cg::this_cluster` to get cluster
- Member functions (highlights)
  - `sync()` Sync in the cluster
  - `thread_rank()` Get rank within cluster
  - `map_shared_rank()` Get address of shared memory of another block of cluster
- See [cluster group documentation](#) and [thread block cluster introduction](#)

# Conclusions

# Conclusions

- **CG** alternative model to create groups
- Groups are **entities**, have member functions
- Synchronizing is important (not mentioned before: `__syncwarps()`)
- **Warp-level functions** easily accessible from groups
- Some new device features only exposed through CG
- See also further literature in [Appendix](#)

Thank you  
for your attention!  
[a.herten@fz-juelich.de](mailto:a.herten@fz-juelich.de)

# Appendix

## Appendix

Further Literature

Glossary

References: Images

# Further Literature

- NVIDIA Developer Blog: [Cooperative Groups: Flexible CUDA Thread Programming](#)
- NVIDIA Developer Blog: [Inside Volta: The World's Most Advanced Data Center GPU](#)
- NVIDIA Developer Blog: [Using CUDA Warp-Level Primitives](#)
- Talk at GPU Technology Conference 2018: [Cooperative Groups](#) by Kyrylo Perelygin and Yuan Lin
- Talk: [Warp-synchronous programming with Cooperative Groups](#) by Sylvain Collange
- Book: [CUDA Programming](#) by Shane Cook

# Glossary I

**API** A programmatic interface to software by well-defined functions. Short for application programming interface. 35, 37, 39, 41

**CUDA** Computing platform for GPUs from NVIDIA. Provides, among others, CUDA C/C++. 4, 30, 42, 43, 45, 48, 50

**NVIDIA** US technology company creating GPUs. 63

**CG** Cooperative Groups. 6, 10, 30, 42, 45, 59

**GPU** Graphics Processing Unit. 63

**SIMD** Single Instruction, Multiple Data. 47

**SIMT** Single Instruction, Multiple Threads. 4, 47

**SM** Streaming Multiprocessor. 48

# References: Images, Graphics I

- [1] Yuriy Rzhemovskiy. *Teenage Penguins*. Freely available at Unsplash. URL: <https://unsplash.com/photos/qFxS5FkUSAQ>.